Spring
IN ACTION

THIRD EDITION

Craig Walls

MANNING

Table of Contents

Working with remote services

This chapter covers

- Accessing and exposing RMI services
- Using Hessian and Burlap services
- Working with Spring's HTTP invoker
- Using Spring with web services

Imagine for a moment that you're stranded on a deserted island. This may sound like a dream come true. After all, who wouldn't want to get some solitude on a beach, blissfully ignorant of the goings-on of the outside world?

But on a deserted island, it's not pina coladas and sunbathing all the time. Even if you enjoy the peaceful seclusion, it won't be long before you'll get hungry, bored, and lonely. You can only live on coconuts and spear-caught fish for so long. You'll eventually need food, fresh clothing, and other supplies. And if you don't get in contact with another human soon, you may end up talking to a volleyball!

Many applications that you'll develop are like island castaways. On the surface they might seem self-sufficient, but in reality, they probably collaborate with other systems, both within your organization and externally.

255

For example, consider a procurement system that needs to communicate with a vendor's supply chain system. Maybe your company's human resources system needs to integrate with the payroll system. Or the payroll system may need to communicate with an external system that prints and mails paychecks. No matter the circumstance, your application will need to communicate with other systems to access services remotely.

Several remoting technologies are available to you, as a Java developer, including

- Remote Method Invocation (RMI)
- Caucho's Hessian and Burlap
- Spring's own HTTP-based remoting
- Web services with JAX-RPC and JAX-WS

Regardless of which remoting technology you choose, Spring provides broad support for accessing and creating remote services with several different technologies. In this chapter, you'll learn how Spring both simplifies and complements these remoting services. But first, let's set the stage for this chapter with an overview of how remoting works in Spring.

10.1 An overview of Spring remoting

Remoting is a conversation between a client application and a service. On the client side, some functionality is required that isn't within the scope of the application. So the application reaches out to another system that can provide the functionality. The remote application exposes the functionality through a remote service.

Suppose that we'd like to make some of the Spitter application's functionality available as remote services for other applications to use. Perhaps in addition to the existing browser-based user interface, we'd like to make a desktop or mobile front end for Spitter, as illustrated in figure 10.1. To support that, we'll need to expose the basic functions of the SpitterService interface as a remote service.

The conversation between the other applications and Spitter begins with a *remote procedure call (RPC)* from the client applications. On the surface, an RPC is similar to a call to a method on a local object. Both are synchronous operations, blocking execution in the calling code until the called procedure is complete.

The difference is a matter of proximity, with an analogy to human communication. If you're at the proverbial watercooler at work discussing the outcome of the weekend's football game, you're conducting a local conversation—the conversation takes place between two people in the same room. Likewise, a local method call is one where execution flow is exchanged between two blocks of code within the same application.


Figure 10.1   A third-party client can interact with the Spitter application by making remote calls to a service exposed by Spitter.

On the other hand, if you were to pick up the phone to call a client in another city, your conversation would be conducted remotely over the telephone network. Similarly, RPC is when execution flow is handed off from one application to another application, theoretically on a different machine in a remote location over the network.

Spring supports remoting for several different RPC models, including Remote Method Invocation (RMI), Caucho's Hessian and Burlap, and Spring's own HTTP invoker. Table 10.1 outlines each of these models and briefly discusses their usefulness in various situations.

Table 10.1   Spring supports RPC via several remoting technologies.

RPC model	Useful when...
Remote Method Invocation (RMI)	Accessing/exposing Java-based services when network constraints such as firewalls aren't a factor
Hessian or Burlap	Accessing/exposing Java-based services over HTTP when network constraints are a factor
HTTP invoker	Accessing/exposing Spring-based services when network constraints are a factor and you desire Java serialization over XML or proprietary serialization
JAX-RPC and JAX-WS	Accessing/exposing platform-neutral, SOAP-based web services

Regardless of which remoting model you choose, you'll find that a common theme runs through Spring's support for each model. This means that once you understand how to configure Spring to work with one of the models, you'll have a modest learning curve if you decide to use a different model.

In all models, services can be configured into your application as Spring-managed beans. This is accomplished using a proxy factory bean that enables you to wire remote services into properties of your other beans as if they were local objects. Figure 10.2 illustrates how this works.

The client makes calls to the proxy as if the proxy were providing the service functionality. The proxy communicates with the remote service on behalf of the client. It handles the details of connecting and making remote calls to the remote service.
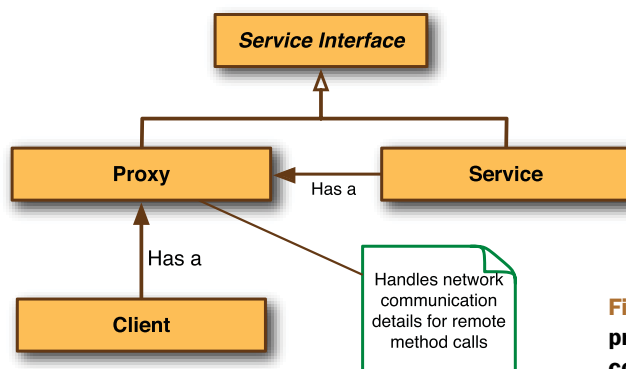

Figure 10.2   In Spring, remote services are proxied so that they can be wired into client code as if they were any other Spring bean.
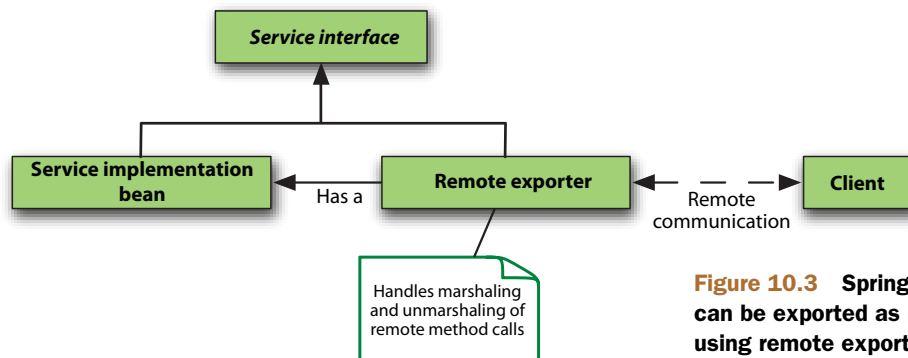
Figure 10.3 Spring-managed beans can be exported as remote services using remote exporters.

What's more, if the call to the remote service results in a `java.rmi.RemoteException`, the proxy handles that exception and rethrows it as an unchecked `RemoteAccess-Exception`. Remote exceptions usually signal problems such as network or configuration issues that can't be gracefully recovered from. Since a client can usually do little to gracefully recover from a remote exception, rethrowing a `RemoteAccessException` makes it optional for the client to handle the exception.

On the service side, you're able to expose the functionality of any Spring-managed bean as a remote service using any of the models listed in table 10.1. Figure 10.3 illustrates how remote exporters expose bean methods as remote services.

Whether you'll be developing code that consumes remote services, implementing those services, or both, working with remote services in Spring is purely a matter of configuration. You won't have to write any Java code to support remoting. Your service beans don't have to be aware that they're involved in an RPC (although any beans passed to or returned from remote calls may need to implement `java.io.Serializable`).

Let's start our exploration of Spring's remoting support by looking at RMI, the original remoting technology for Java.

10.2 *Working with RMI*

If you've been working in Java for any length of time, you've no doubt heard of (and probably used) Remote Method Invocation (RMI). RMI—first introduced into the Java platform in JDK 1.1—gives Java programmers a powerful way to communicate between Java programs. Before RMI, the only remoting options available to Java programmers were CORBA (which at the time required the purchase of a third-party *Object Request Broker*, or *ORB*) or handwritten socket programming.

But developing and accessing RMI services is tedious, involving several steps, both programmatic and manual. Spring simplifies the RMI model by providing a proxy factory bean that enables you to wire RMI services into your Spring application is if they were local JavaBeans. Spring also provides a remote exporter that makes short work of converting your Spring-managed beans into RMI services.

For the Spitter application, I'll show you how to wire an RMI service into a client application's Spring application context. But first, let's see how to use the RMI exporter to publish the `SpitterService` implementation as an RMI service.

10.2.1 *Exporting an RMI service*

If you've ever created an RMI service before, you know that it involves the following steps:

1 Write the service implementation class with methods that throw `java.rmi` `.RemoteException`.
2 Create the service interface to extend `java.rmi.Remote`.
3 Run the RMI compiler (`rmic`) to produce client stub and server skeleton classes.
4 Start an RMI registry to host the services.
5 Register the service in the RMI registry.

Wow! That's a lot of work just to publish a simple RMI service. What's perhaps worse than all the steps required, you may have noticed that `RemoteExceptions` and `MalformedURLExceptions` are thrown around a lot. These exceptions usually indicate a fatal error that can't be recovered from in a catch block, but you're still expected to write boilerplate code that catches and handles those exceptions—even if there's not much you can do to fix them.

Clearly a lot of code and manual work is involved to publish an RMI service. Is there anything Spring can do to make this situation less knotty?

CONFIGURING AN RMI SERVICE IN SPRING

Fortunately, Spring provides an easier way to publish RMI services. Instead of writing RMI-specific classes with methods that throw `RemoteException`, you simply write a POJO that performs the functionality of your service. Spring handles the rest.

The RMI service that we'll create exposes the methods from the `SpitterService` interface. As a reminder, the following shows what that interface looks like.

> **Listing 10.1 `SpitterService` defines the service layer of the Spitter application.**

```
package com.habuma.spitter.service;

import java.util.List;

import com.habuma.spitter.domain.Spitter;
import com.habuma.spitter.domain.Spittle;

public interface SpitterService {
  List<Spittle> getRecentSpittles(int count);
  void saveSpittle(Spittle spittle);

  void saveSpitter(Spitter spitter);
  Spitter getSpitter(long id);
  void startFollowing(Spitter follower, Spitter followee);

  List<Spittle> getSpittlesForSpitter(Spitter spitter);
  List<Spittle> getSpittlesForSpitter(String username);
  Spitter getSpitter(String username);

  Spittle getSpittleById(long id);
  void deleteSpittle(long id);

  List<Spitter> getAllSpitters();
}
```

If we were using traditional RMI to expose this service, all of those methods in `SpitterService` and in `SpitterServiceImpl` would need to throw `java.rmi.Remote-Exception`. But we're going to turn it into an RMI service using Spring's `RmiService-Exporter`, so the existing implementations will do fine.

`RmiServiceExporter` exports any Spring-managed bean as an RMI service. As depicted in figure 10.4, `RmiServiceExporter` works by wrapping the bean in an adapter class. The adapter class is then bound to the RMI registry and proxies requests to the service class—`SpitterServiceImpl` in this case.

The simplest way to use `RmiServiceExporter` to expose `SpitterServiceImpl` as an RMI service is to configure it in Spring with the following XML:

```
<bean class="org.springframework.remoting.rmi.RmiServiceExporter"
      p:service-ref="spitterService"
      p:serviceName="SpitterService"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Here the `spitterService` bean is wired into the `service` property to indicate that the `RmiServiceExporter` is to export the bean as an RMI service. The `serviceName` property names the RMI service. And the `serviceInterface` property specifies the interface that the service implements.

By default `RmiServiceExporter` attempts to bind to an RMI registry on port 1099 of the local machine. If no RMI registry is found at that port, `RmiServiceExporter` will start one. If you'd rather bind to an RMI registry at a different port or host, you can specify so with the `registryPort` and `registryHost` properties. For example, the following `RmiServiceExporter` will attempt to bind to an RMI registry on port 1199 on the host rmi.spitter.com:

```
<bean class="org.springframework.remoting.rmi.RmiServiceExporter"
      p:service-ref="spitterService"
      p:serviceName="SpitterService"
      p:serviceInterface="com.habuma.spitter.service.SpitterService"
      p:registryHost="rmi.spitter.com"
      p:registryPort="1199"/>
```

That's all you need to do to have Spring turn a bean into an RMI service. Now that the Spitter service has been exposed as an RMI service, we may create alternative user interfaces or invite third parties to create new clients for Spitter that use the RMI service. The developers of those clients will have an easy time connecting to the Spitter


Figure 10.4 `RmiServiceExporter` turns POJOs into RMI services by wrapping them in a service adapter and binding the service adapter to the RMI registry.

RMI service if they're using Spring. Let's switch gears and see how to write a client of the Spitter RMI service.

10.2.2 *Wiring an RMI service*

Traditionally, RMI clients must use the RMI API's `Naming` class to look up a service from the RMI registry. For example, the following snippet of code might be used to retrieve the RMI Spitter service:

```
try {
  String serviceUrl = "rmi:/spitter/SpitterService";
  SpitterService spitterService =
          (SpitterService) Naming.lookup(serviceUrl);
  ...
}
catch (RemoteException e) { ... }
catch (NotBoundException e) { ... }
catch (MalformedURLException e) { ... }
```

Though this snippet of code would certainly retrieve a reference to the RMI spitter service, it presents two problems:

- Conventional RMI lookups could result in any one of three checked exceptions (`RemoteException`, `NotBoundException`, and `MalformedURLException`) that must be caught or rethrown.
- Any code that needs the spitter service is responsible for retrieving the service itself. That's plumbing code and is probably not directly cohesive with the client's functionality.

The exceptions thrown in the course of an RMI lookup are the kinds that typically signal a fatal and unrecoverable condition in the application. `MalformedURLException`, for instance, indicates that the address given for the service isn't valid. To recover from this exception, the application will at a minimum need to be reconfigured and may have to be recompiled. No try/catch block will be able to recover gracefully, so why should your code be forced to catch and handle it?

But perhaps more sinister is the fact that this code is in direct opposition to the principles of dependency injection. Because the client code is responsible for looking up the Spitter service *and* that the service is an RMI service, there's no opportunity to provide a different implementation of `SpitterService` from some other source. Ideally, you should be able to inject a `SpitterService` object into any bean that needs one instead of having the bean look up the service itself. Using DI, any client of `SpitterService` can be ignorant of where that service comes from.

Spring's `RmiProxyFactoryBean` is a factory bean that creates a proxy to an RMI service. Using `RmiProxyFactoryBean` to reference an RMI `SpitterService` is as simple as adding the following <bean> to the client's Spring configuration file:

```
<bean id="spitterService"
      class="org.springframework.remoting.rmi.RmiProxyFactoryBean"
      p:serviceUrl="rmi://localhost/SpitterService"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

The URL of the service is set through `RmiProxyFactoryBean`'s `serviceUrl` property. Here, the service is named `SpitterService` and is hosted on the local machine. Meanwhile, the interface that the service provides is specified with the `service-Interface` property. The interaction between the client and the RMI proxy is illustrated in figure 10.5.

Now that we've declared the RMI service as a Spring-managed bean, we can wire it as a dependency into another bean just as you would a regular non-remote bean. For example, suppose the client needs to use the Spitter service to retrieve a list of `Spittle`s for a given user. You might use `@Autowired` to wire the service proxy into the client:

```
@Autowired
  SpitterService spitterService;
```

Then you can invoke methods on it as if it were a local bean:

```
public List<Spittle> getSpittles(String userName) {
  Spitter spitter = spitterService.getSpitter(userName);
  return spitterService.getSpittlesForSpitter(spitter);
}
```

What's great about accessing an RMI service in this way is that the client code doesn't even know that it's dealing with an RMI service. It's given a `SpitterService` object via injection, without any concern for where it comes from.

Furthermore, the proxy catches any `RemoteException`s that may be thrown by the service and rethrows them as unchecked exceptions that you may safely ignore. This makes it possible to easily swap out the remote service bean with another implementation of the service—perhaps a different remote service or maybe a mock implementation used when unit testing the client code.

Even though the client code isn't aware that the `SpitterService` that it was given is a remote service, you may want to take care designing the service's interface. Note
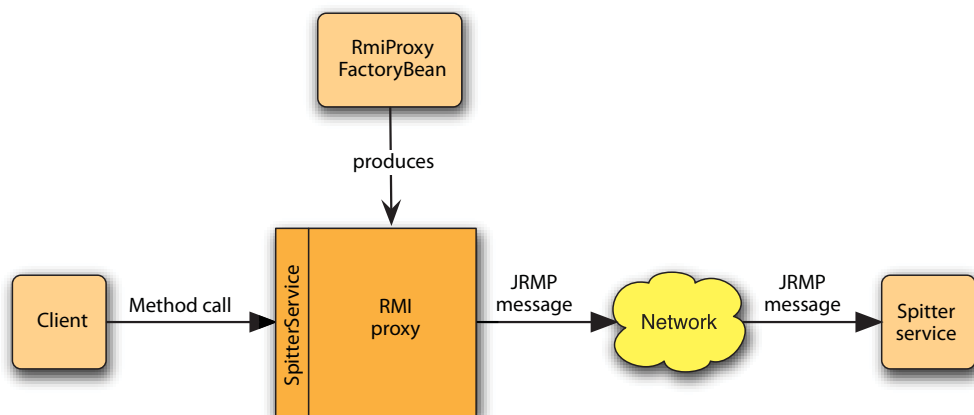

Figure 10.5   `RmiProxyFactoryBean` **produces a proxy object that talks to remote RMI services on behalf of the client. The client talks to the proxy through the service's interface as if the remote service were just a local POJO.**

that the client had to make two calls to the service: one to look up the `Spitter` by their username and another to retrieve the list of `Spittle` objects. That's two remote calls that are affected by network latency and will impact the performance of the client. Knowing that this is how the service will be used, it may be worthwhile to revisit the service's interface to consolidate those two calls into a single method. But for now we'll accept the service as is.

RMI is an excellent way to communicate with remote services, but it has some limitations. First, RMI has difficulty working across firewalls. That's because RMI uses arbitrary ports for communication—something that firewalls typically won't allow. In an intranet environment, this usually isn't a concern. But if you're working on the "evil internet," you'll probably run into trouble with RMI. Even through RMI has support for tunneling over HTTP (which is usually allowed by firewalls), setting up RMI tunneling can be tricky.

Another thing to consider is that RMI is Java-based. That means that both the client and the service must be written in Java. And since RMI uses Java serialization, the types of the objects being sent across the network must have the exact same version on both sides of the call. These may or may not be issues for your application, but bear them in mind when choosing RMI for remoting.

Caucho Technology (the same people behind the Resin application server) has developed a remoting solution that addresses the limitations of RMI. Actually, they've come up with two solutions: Hessian and Burlap. Let's see how to use Hessian and Burlap to work with remote services in Spring.

10.3   Exposing remote services with Hessian and Burlap

Hessian and Burlap are two solutions provided by Caucho Technology[1] that enable lightweight remote services over HTTP. They each aim to simplify web services by keeping both their API and their communication protocols as simple as possible.

You may be wondering why Caucho has two solutions to the same problem. Hessian and Burlap are two sides of the same coin, but each serves slightly different purposes. Hessian, like RMI, uses binary messages to communicate between client and service. But unlike other binary remoting technologies (such as RMI), the binary message is portable to languages other than Java, including PHP, Python, C++, and C#.

Burlap is an XML-based remoting technology, which automatically makes it portable to any language that can parse XML. And because it's XML, it's more easily human-readable than Hessian's binary format. Unlike other XML-based remoting technologies (such as SOAP or XML-RPC), Burlap's message structure is as simple as possible and doesn't require an external definition language (such as WSDL or IDL).

You may be wondering how to choose between Hessian and Burlap. For the most part, they're identical. The only difference is that Hessian messages are binary and Burlap messages are XML. Because Hessian messages are binary, they're more

---

[1] http://www.caucho.com

bandwidth friendly. If human readability is important to you (for debugging purposes) or if your application will be communicating with a language for which there's no Hessian implementation, Burlap's XML messages may be preferable.

To demonstrate Hessian and Burlap services in Spring, let's revisit the Spitter service example that we addressed with RMI in the previous section. But this time we'll look at how to solve the problem using Hessian and Burlap as the remoting models.

10.3.1 *Exposing bean functionality with Hessian/Burlap*

As before, suppose that we want to expose the functionality of the `SpitterService-Impl` class as a service—a Hessian service this time around. Even without Spring, this would be fairly trivial. You'd write a service class that extends `com.caucho.hessian.server.HessianServlet` and make sure that all of the service methods are `public` (all public methods are considered service methods to Hessian).

Because Hessian services are already easy to implement, Spring doesn't do much to simplify the Hessian model any further. But when used with Spring, a Hessian service can take full advantage of the Spring Framework in ways that a pure Hessian service can't. This includes using Spring AOP to advise a Hessian service with system-wide services such as declarative transactions.

EXPORTING A HESSIAN SERVICE

Exporting a Hessian service in Spring is remarkably similar to implementing an RMI service in Spring. To expose the Spitter service bean as an RMI service, we had to configure an `RmiServiceExporter` bean in the Spring configuration file. In a similar way, to expose the Spitter service as a Hessian service, we'll need to configure another exporter bean. This time it'll be a `HessianServiceExporter`.

`HessianServiceExporter` performs the same function for a Hessian service as `RmiServiceExporter` does for an RMI service: it exposes the public methods of a POJO as methods of a Hessian service. But, as shown in figure 10.6, how it pulls off this feat is different from how `RmiServiceExporter` exports POJOs as RMI services.

`HessianServiceExporter` is a Spring MVC controller (more on that in a moment) that receives Hessian requests and translates them into method calls on the exported POJO.


Figure 10.6 `HessianService-Exporter` **is a Spring MVC controller that exports a POJO as a Hessian service by receiving Hessian requests and translating them into calls to the POJO.**

The following declaration of `HessianServiceExporter` in Spring exports the `spitterService` bean as a Hessian service:

```
<bean id="hessianSpitterService"
      class="org.springframework.remoting.caucho.HessianServiceExporter"
      p:service-ref="spitterService"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Just as with `RmiServiceExporter`, the `service` property is wired with a reference to the bean that implements the service. Here, that's a reference to the `spitterService` bean. The `serviceInterface` is set to indicate that `SpitterService` is the interface that the service implements.

Unlike `RmiServiceExporter`, we don't need to get a `serviceName` property. With RMI, the `serviceName` property is used to register a service in the RMI registry. Hessian doesn't have a registry and, therefore, there's no need to name a Hessian service.

CONFIGURING THE HESSIAN CONTROLLER

Another major difference between `RmiServiceExporter` and `HessianServiceExporter` is that because Hessian is HTTP-based, `HessianServiceExporter` is implemented as a Spring MVC Controller. This means that in order to use exported Hessian services, you'll need to perform two additional configuration steps:

- Configure a Spring `DispatcherServlet` in web.xml and deploy your application as a web application.
- Configure a URL handler in your Spring configuration file to dispatch Hessian service URLs to the appropriate Hessian service bean.

We saw how to configure Spring's `DispatcherServlet` and URL handlers in chapter 7. So these steps should be somewhat familiar by now. First, you'll need a `DispatcherServlet`. Fortunately, you have one already configured in the Spitter application's web.xml file. But for the purposes of handling Hessian services, that `DispatcherServlet` will need a servlet mapping that catches `*.service` URLs:

```
<servlet-mapping>
  <servlet-name>spitter</servlet-name>
  <url-pattern>*.service</url-pattern>
</servlet-mapping>
```

Configured this way, any request whose URL ends with `.service` will be given to `DispatcherServlet`, which will in turn hand off the request to the `Controller` that's mapped to the URL. Thus requests to `/spitter.service` will ultimately be handled by the `hessianSpitterService` bean (which is just a proxy to `SpitterServiceImpl`).

How do I know that the request will go to `hessianSpitterService`? Because we're also going to configure a URL mapping to have `DispatcherServlet` send it to `hessianSpitterService`. The following `SimpleUrlHandlerMapping` will make that happen:

```
<bean id="urlMapping" class=
    "org.springframework.web.servlet.handler.SimpleUrlHandlerMapping">
  <property name="mappings">
    <value>
```

```
        /spitter.service=hessianSpitterService
    </value>
  </property>
</bean>
```

An alternative to Hessian's binary protocol is Burlap's XML-based protocol. Let's see how to export a service as a Burlap service.

EXPORTING A BURLAP SERVICE

`BurlapServiceExorter` is virtually identical to `HessianServiceExporter` in every way, except that it uses an XML-based protocol instead of a binary protocol. The following bean definition shows how to expose the Spitter service as a Burlap service using `BurlapServiceExporter`:

```
<bean id="burlapSpitterService"
      class="org.springframework.remoting.caucho.BurlapServiceExporter"
      p:service-ref="spitterService"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

As you can see, the only thing different between this bean and its Hessian counterpart are the bean's ID and its class. Configuring a Burlap service is otherwise the same as configuring a Hessian service. This includes the need to set up a URL handler and a `DispatcherServlet`.

Now let's look at the other side of the conversation and consume the service that we published using Hessian (or Burlap).

10.3.2 *Accessing Hessian/Burlap services*

As you'll recall from section 10.2.2, client code that consumed the Spitter service using `RmiProxyFactoryBean` had no idea that the service was an RMI service. In fact, it had no clue that it was a remote service at all. It only dealt with the `SpitterService` interface—all of the RMI details were completely contained in the configuration of the beans in Spring's configuration file. The good news is that because of the client's ignorance of the service's implementation, switching from an RMI client to a Hessian client is extremely easy, requiring no changes to the client's Java code.

The bad news is that if you love writing Java code, this section may be a letdown. That's because the only difference between wiring the client side of an RMI-based service and wiring the client side of a Hessian-based service is that you'll use Spring's `HessianProxyFactoryBean` instead of `RmiProxyFactoryBean`. A Hessian-based spitter service can be declared in the client code like this:

```
<bean id="spitterService"
      class="org.springframework.remoting.caucho.HessianProxyFactoryBean"
      p:serviceUrl="http://localhost:8080/Spitter/spitter.service"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Just as with an RMI-based service, the `serviceInterface` property specifies the interface that the service implements. And, as with `RmiProxyFactoryBean`, `serviceUrl` indicates the URL of the service. Since Hessian is HTTP-based, it's been set to an HTTP
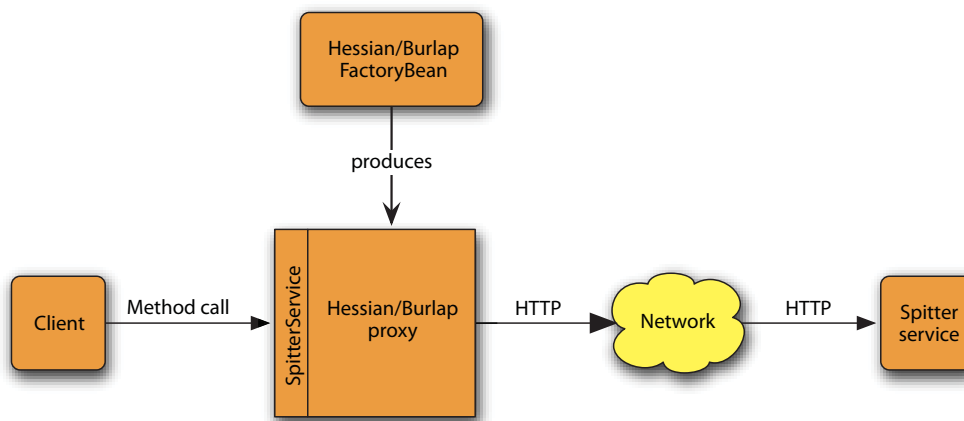
Figure 10.7 `HessianProxyFactoryBean` and `BurlapProxyFactoryBean` produce proxy objects that talk to a remote service over HTTP (Hessian in binary, Burlap in XML).

URL here (determined in part by the URL mapping we defined earlier). Figure 10.7 shows the interaction between a client and the proxy produced by `HessianProxy-FactoryBean`.

As it turns out, wiring a Burlap service into the client is equally uninteresting. The only difference is that you'll use `BurlapProxyFactoryBean` instead of `HessianProxy-FactoryBean`:

```
<bean id="spitterService"
      class="org.springframework.remoting.caucho.BurlapProxyFactoryBean"
      p:serviceUrl="http://localhost:8080/Spitter/spitter.service"
      p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Although I've made light of how uninteresting the configuration differences are among RMI, Hessian, and Burlap, this tedium is a benefit. It demonstrates that you can switch effortlessly between the various remoting technologies supported by Spring without having to learn a completely new model. Once you've configured a reference to an RMI service, it's short work to reconfigure it as a Hessian or Burlap service.

Because both Hessian and Burlap are based on HTTP, they don't suffer from the same firewall issues as RMI. But RMI has both Hessian and Burlap beat when it comes to serializing objects that are sent in RPC messages. Whereas Hessian and Burlap both use a proprietary serialization mechanism, RMI uses Java's own serialization mechanism. If your data model is complex, the Hessian/Burlap serialization model may not be sufficient.

There is a best-of-both-worlds solution. Let's look at Spring's HTTP invoker, which offers RPC over HTTP (like Hessian/Burlap) while at the same time using Java serialization of objects (like RMI).

10.4   Using Spring's HttpInvoker

The Spring team recognized a void between RMI services and HTTP-based services such as Hessian and Burlap. On one side, RMI uses Java's standard object serialization but is difficult to use across firewalls. On the other side, Hessian and Burlap work well across firewalls but use a proprietary object serialization mechanism.

Thus Spring's HTTP invoker was born. The HTTP invoker is a new remoting model created as part of the Spring Framework to perform remoting across HTTP (to make the firewalls happy) and using Java's serialization (to make programmers happy).

Working with HTTP invoker-based services is similar to working with Hessian/Burlap-based services. To get started with the HTTP invoker, let's take another look at the Spitter service—this time implemented as an HTTP invoker service.

10.4.1   Exposing beans as HTTP services

To export a bean as an RMI service, we used `RmiServiceExporter`. To export it as a Hessian service, we used `HessianServiceExporter`. And to export it as a Burlap service, we used `BurlapServiceExporter`. Continuing this monotony over to Spring's HTTP invoker, it shouldn't surprise you that to export an HTTP invoker service, we'll need to use `HttpInvokerServiceExporter`.

To export the Spitter service as an HTTP invoker–based service, you need to configure an `HttpInvokerServiceExporter` bean like this:

```
<bean class=
     "org.springframework.remoting.httpinvoker.HttpInvokerServiceExporter"
     p:service-ref="spitterService"
     p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Feeling a sense of déjà vu? You may have a hard time spotting the difference between this bean declaration and the ones in section 10.3.2. The only material difference is the class name: `HttpInvokerServiceExporter`. Otherwise, this exporter isn't much different from the other remote service exporters.

As illustrated in figure 10.8, `HttpInvokerServiceExporter` works much like `HessianServiceExporter` and `BurlapServiceExporter`. It's a Spring MVC controller that receives requests from a client through `DispatcherServlet` and translates those requests into method calls on the service implementation POJO.
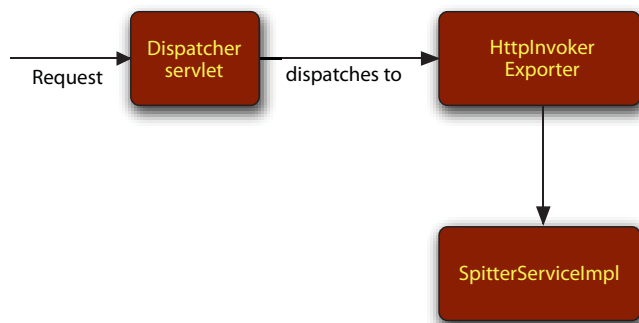

Figure 10.8 `HttpInvokerServiceExporter` works much like its Hessian and Burlap cousins, receiving requests from a Spring MVC `DispatcherServlet` and translating them into method calls on a Spring-managed bean.

Because `HttpInvokerServiceExporter` is a Spring MVC controller, you'll need to set up a URL handler to map an HTTP URL to the service, just like with the Hessian and Burlap exporters:

```
<bean id="urlMapping" class=
    "org.springframework.web.servlet.handler.SimpleUrlHandlerMapping">
  <property name="mappings">
    <value>
      /spitter.service=httpInvokerSpitterService
    </value>
  </property>
</bean>
```

Also, as before, you'll need to make sure you have a `DispatcherServlet` declared in web.xml with the following `<servlet-mapping>`:

```
<servlet-mapping>
  <servlet-name>spitter</servlet-name>
  <url-pattern>*.service</url-pattern>
</servlet-mapping>
```

Configured this way, the Spitter service will be available at /spitter.service, the same URL you used to expose the service through Hessian and Burlap.

We've already seen how to consume services made remote through RMI, Hessian, and Burlap. Now let's rework the Spitter client to use the service that you just exposed with HTTP invoker.

10.4.2 Accessing services via HTTP

At the risk of sounding like a broken record, I must tell you that consuming an HTTP invoker-based service is much like what we've already seen with the other remote service proxies. It's virtually identical. As you can see from figure 10.9, `HttpInvoker-ProxyFactoryBean` fills the same hole as the other remote service proxy factory beans that we've seen in this chapter.
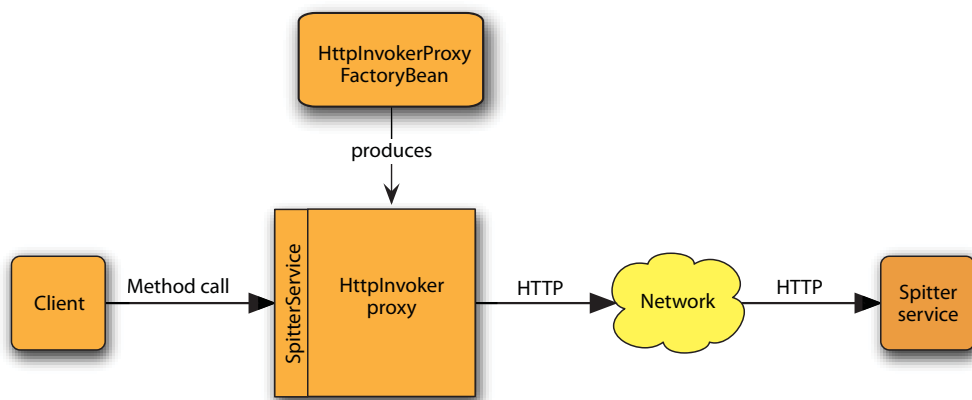

Figure 10.9 `HttpInvokerProxyFactoryBean` is a proxy factory bean that produces a proxy for remoting with a Spring-specific HTTP-based protocol.

To wire the HTTP invoker–based service into our client's Spring application context, we must configure a bean that proxies it using `HttpInvokerProxyFactoryBean` as follows:

```
<bean id="spitterService" class=
    "org.springframework.remoting.httpinvoker.HttpInvokerProxyFactoryBean"
    p:serviceUrl="http://localhost:8080/Spitter/spitter.service"
    p:serviceInterface="com.habuma.spitter.service.SpitterService" />
```

Comparing this bean definition to those in sections 10.2.2 and 10.3.2, you'll find that little has changed. The `serviceInterface` property is still used to indicate the interface implemented by the Spitter service. And the `serviceUrl` property is still used to indicate the location of the remote service. Because HTTP invoker is HTTP-based like Hessian and Burlap, the `serviceUrl` can contain the same URL as with the Hessian and Burlap versions of the bean.

Don't you love the symmetry?

Spring's HTTP invoker presents a best-of-both-worlds remoting solution combining the simplicity of HTTP communication with Java's built-in object serialization. This makes HTTP invoker services an appealing alternative to either RMI or Hessian/Burlap.

`HttpInvoker` has one significant limitation that you should keep in mind: it's a remoting solution offered by the Spring Framework only. This means that both the client and the service must be Spring-enabled applications. This also implies, at least for now, that both the client and the service must be Java-based. And because Java serialization is being used, both sides must have the same version of the classes (much like RMI).

RMI, Hessian, Burlap, and HTTP invoker are great remoting options. But when it comes to ubiquitous remoting, none hold a candle to web services. Next up, we'll look at how Spring supports remoting through SOAP-based web services.

10.5 Publishing and consuming web services

One of the most hyped TLAs (three-letter acronyms) in recent years is SOA (which stands for *service-oriented architecture*). SOA means many things to different people. But at the center of SOA is the idea that applications can and should be designed to lean on a common set of core services instead of reimplementing the same functionality for each application.

For example, a financial institution may have several applications, many of which need access to borrower account information. Rather than build account access logic into each application (much of which would be duplicated), the applications could all rely on a common service to retrieve the account information.

Java and web services have a long history together and several options are available for working with web services in Java. Many of those options integrate with Spring in some way. Though it'd be impossible for me to cover every Spring-enabled web service framework and toolkit in this book, Spring itself comes with some capable support for publishing and consuming SOAP web services using the Java API for XML Web Services, or JAX-WS as it's commonly known.

> ### What about JAX-RPC and XFire?
>
> In previous editions of this book, I wrote about developing web services using XFire (http://xfire.codehaus.org) and Spring's support for JAX-RPC. Those were great topics back then, but both are fading technologies.
>
> JAX-RPC was supplanted by JAX-WS as the web service standard for Java. Spring followed suit and has now deprecated its support for JAX-RPC in favor of the new JAX-WS support. Fortunately, the JAX-WS support in Spring closely mirrors the JAX-RPC support. Spring's `JaxWsPortProxyFactoryBean`, for instance, works much like the old `JaxRpcPortProxyFactoryBean`.
>
> XFire was my favorite way of working with web services in Spring. But development of XFire stopped with the 1.2.6 release. The Apache CXF (http://cxf.apache.org) project is considered by many to be XFire 2; so if you liked XFire, then you should check out Apache CXF. Apache CXF is far more ambitious than XFire and to cover it adequately would be well outside of the scope of this book.
>
> Since one of my objectives for this edition was to remain as current as possible, I chose to leave JAX-RPC and XFire behind. If either of those topics interest you, I encourage you to find a copy of *Spring in Action, Second Edition*. Both topics are covered there and little has changed with regard to JAX-RPC or XFire since that time.

In this section, we'll revisit the Spitter service example one more time. This time, we'll expose and consume the Spitter service as a web service using Spring's JAX-WS support. Let's start by seeing what it takes to create a JAX-WS web service in Spring.

10.5.1 Creating Spring-enabled JAX-WS endpoints

Earlier in this chapter, we created remote services using Spring's service exporters. These service exporters magically turn Spring-configured POJOs into remote services. We saw how to create RMI services using `RmiServiceExporter`, Hessian services using `HessianServiceExporter`, Burlap services using `BurlapServiceExporter`, and HTTP invoker services using `HttpInvokerServiceExporter`. Now you probably expect me to show you how to create web services using a JAX-WS service exporter in this section.

Spring does provide a JAX-WS service exporter, `SimpleJaxWsServiceExporter`, and we'll see it soon enough. But before we get there, you should know that it may not be the best choice in all situations. You see, `SimpleJaxWsServiceExporter` requires that the JAX-WS runtime support publishing of endpoints to a specified address.[2] The JAX-WS runtime that ships with Sun's JDK 1.6 fits the bill, but other JAX-WS implementations, including the reference implementation of JAX-WS, may not.

If you'll be deploying to a JAX-WS runtime that doesn't support publishing to a specified address, then you'll have write your JAX-WS endpoints in a more conventional way. That means that the lifecycle of the endpoints will be managed by the JAX-

---

[2] More specifically, this means that the JAX-WS provider must come with its own HTTP server that it can use to build up the necessary infrastructure for publishing the service to a requested address.

WS runtime and not by Spring. But that doesn't mean that they can't be wired with
beans from a Spring application context.

AUTOWIRING JAX-WS ENDPOINTS IN SPRING

The JAX-WS programming model involves using annotations to declare a class and its
methods as web service operations. A class that's annotated with `@WebService` is con-
sidered a web service endpoint and its methods, annotated with `@WebMethod`, are the
operations.

Just as with any other object in a sizable application, a JAX-WS endpoint will likely
depend on other objects to do its work. That means that JAX-WS endpoints could ben-
efit from dependency injection. But if the endpoint's lifecycle is managed by the JAX-
WS runtime and not by Spring, it would seem to be impossible to wire Spring-
managed beans into a JAX-WS managed endpoint instance.

The secret to wiring JAX-WS endpoints is to extend `SpringBeanAutowiring-`
`Support`. By extending `SpringBeanAutowiringSupport`, you can annotate an end-
point's properties with `@Autowired` and its dependencies will be met.[3] `Spitter-`
`ServiceEndpoint` shows how this works.

Listing 10.2 `SpringBeanAutowiringSupport` on JAX-WS endpoints

```
package com.habuma.spitter.remoting.jaxws;
import java.util.List;

import javax.jws.WebMethod;
import javax.jws.WebService;

import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.web.context.support.SpringBeanAutowiringSupport;

import com.habuma.spitter.domain.Spitter;
import com.habuma.spitter.domain.Spittle;
import com.habuma.spitter.service.SpitterService;

@WebService(serviceName="SpitterService")
public class SpitterServiceEndpoint                        Enable
    extends SpringBeanAutowiringSupport {                  autowiring

  @Autowired                                               Autowire
  SpitterService spitterService;                           SpitterService

  @WebMethod
  public void addSpittle(Spittle spittle) {
    spitterService.saveSpittle(spittle);
  }
                                                           Delegate to
  @WebMethod                                               SpitterService
  public void deleteSpittle(long spittleId) {
    spitterService.deleteSpittle(spittleId);
```

---

[3] Although we're using `SpringBeanAutowiringSupport` to enable autowiring for JAX-WS endpoints, it's use-
ful to enable autowiring support anywhere an object's lifecycle is managed outside of Spring. The only
requirement is that the Spring application context and the non-Spring runtime reside in the same web
application.

```
  }

  @WebMethod
  public List<Spittle> getRecentSpittles(int spittleCount) {
    return spitterService.getRecentSpittles(spittleCount);
  }

  @WebMethod
  public List<Spittle> getSpittlesForSpitter(Spitter spitter) {
    return spitterService.getSpittlesForSpitter(spitter);
  }
}
```

Delegate to SpitterService

We've annotated the `spitterService` property with `@Autowired` to indicate that it should be automatically injected with a bean from the Spring application context. From there, this endpoint delegates to the injected `SpitterService` to do the real work.

EXPORTING STANDALONE JAX-WS ENDPOINTS

As I said, `SpringBeanAutowiringSupport` is useful when the object whose properties are being injected doesn't have its lifecycle managed by Spring. But under the right circumstances, it's possible to export a Spring-managed bean as a JAX-WS endpoint.

Spring's `SimpleJaxWsServiceExporter` works in a way similar to the other service exporters that we saw earlier in this chapter, in that it publishes Spring-managed beans as service endpoints in a JAX-WS runtime. Unlike those other service exporters, `SimpleJaxWsServiceExporter` doesn't need to be given a reference to the bean it's supposed to export. Instead, it publishes all beans that are annotated with JAX-WS annotations as JAX-WS services.

`SimpleJaxWsServiceExporter` can be configured using the following `<bean>` declaration:

```
<bean class=
    "org.springframework.remoting.jaxws.SimpleJaxWsServiceExporter"/>
```

As you can see, `SimpleJaxWsServiceExporter` needs nothing else to do its job. When it gets started, it'll dig through the Spring application context looking for beans that are annotated with `@WebService`. When it finds one, it'll publish it as a JAX-WS endpoint with a base address of http://localhost:8080/.

One such bean that it may find is `SpitterServiceEndpoint`.

> **Listing 10.3** `SimpleJaxWsServiceExporter` **turns beans into JAX-WS endpoints.**

```
package com.habuma.spitter.remoting.jaxws;
import java.util.List;

import javax.jws.WebMethod;
import javax.jws.WebService;

import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.stereotype.Component;

import com.habuma.spitter.domain.Spitter;
import com.habuma.spitter.domain.Spittle;
```

```
import com.habuma.spitter.service.SpitterService;

@Component
@WebService(serviceName="SpitterService")
public class SpitterServiceEndpoint {
  @Autowired
  SpitterService spitterService;              ◁──  Autowire
                                                     SpitterService
  @WebMethod
  public void addSpittle(Spittle spittle) {                    ◁┐
    spitterService.saveSpittle(spittle);                        │
  }                                                             │
                                              Delegate to       │
  @WebMethod                                  SpitterService    │
  public void deleteSpittle(long spittleId) {                  │
    spitterService.deleteSpittle(spittleId);                 ◁─┤
  }                                                             │
                                                                │
  @WebMethod                                                    │
  public List<Spittle> getRecentSpittles(int spittleCount) {   │
    return spitterService.getRecentSpittles(spittleCount);   ◁─┤
  }                                                             │
                                                                │
  @WebMethod                                                    │
  public List<Spittle> getSpittlesForSpitter(Spitter spitter) {│
    return spitterService.getSpittlesForSpitter(spitter);    ◁─┘
  }
}
```

You'll notice that this new implementation of `SpitterServiceEndpoint` no longer extends `SpringBeanAutowiringSupport`. As a full-fledged Spring bean, it'll qualify for autowiring without extending any special support class.

   Since `SimpleJaxWsServiceEndpoint`'s base address defaults to http://localhost:8080/ and since `SpitterServiceEndpoint` is annotated with `@WebService(serviceName="SpitterService")`, the matchup of these two beans will result in a web service at http://localhost:8080/SpitterService. But you're in total control of the service URL, so if you'd like, you can set the base address to something else. For example, the following configuration of `SimpleJaxWsServiceEndpoint` publishes the same service endpoint to http://localhost:8888/services/SpitterService.

```
<bean class=
  "org.springframework.remoting.jaxws.SimpleJaxWsServiceExporter"
     p:baseAddress="http://localhost:8888/services/"/>
```

As simple as `SimpleJaxWsServiceEndpoint` seems, you should be aware that it only works with a JAX-WS runtime that supports publication of endpoints with an address. That includes the JAX-WS runtime that comes with Sun's 1.6 JDK. Other JAX-WS runtimes, such as the JAX-WS 2.1 reference implementation, don't support this type of endpoint publication and thus can't be used with `SimpleJaxWsServiceEndpoint`.

10.5.2   Proxying JAX-WS services on the client side

Publishing web services with Spring turned out to be quite different from how we published services in RMI, Hessian, Burlap, and Http invoker. But as you'll soon see,
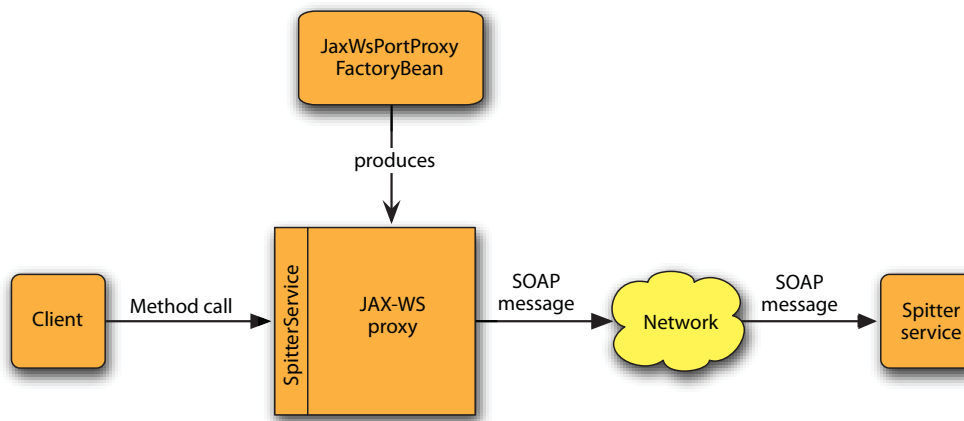
Figure 10.10 `JaxWsPortProxyFactoryBean` **produces proxies that talk to remote web services. These proxies can then be wired into other beans as if they were local POJOs.**

consuming web services with Spring involves client-side proxies in much the same way that Spring-based clients consume those other remoting technologies.

Using `JaxWsPortProxyFactoryBean`, we can wire the Spitter web service in Spring as if it were any other bean. `JaxWsPortProxyFactoryBean` is a Spring FactoryBean that produces a proxy that knows how to talk to a SOAP web service. The proxy itself is created to implement the service's interface (see figure 10.10). Consequently, `Jax-WsPortProxyFactoryBean` makes it possible to wire and use a remote web service as if it were just any other local POJO.

We'll configure `JaxWsPortProxyFactoryBean` to reference the Spitter service web service like this:

```
<bean id="spitterService"
    class="org.springframework.remoting.jaxws.JaxWsPortProxyFactoryBean"
    p:wsdlDocumentUrl="http://localhost:8080/services/SpitterService?wsdl"
    p:serviceName="spitterService"
    p:portName="spitterServiceHttpPort"
    p:serviceInterface="com.habuma.spitter.service.SpitterService"
    p:namespaceUri="http://spitter.com"/>
```

As you can see, several properties must be set for `JaxWsPortProxyFactoryBean` to work. The `wsdlDocumentUrl` property identifies the location of the remote web service's definition file. `JaxWsPortProxyFactoryBean` will use the WSDL available at that URL to construct a proxy to the service. The proxy that's produced by `JaxWsPort-ProxyFactoryBean` will implement the `SpitterService` interface, as specified by the `serviceInterface` property.

The values for the remaining three properties can usually be determined by looking at the service's WSDL. For illustration's sake, suppose that the WSDL for the Spitter service looked like this:

```
<wsdl:definitions targetNamespace="http://spitter.com">
...
  <wsdl:service name="spitterService">
```

```
    <wsdl:port name="spitterServiceHttpPort"
            binding="tns:spitterServiceHttpBinding">
...
    </wsdl:port>
  </wsdl:service>
</wsdl:definitions>
```

Though not likely, it's possible for multiple services and/or ports to be defined in the service's WSDL. For that reason, `JaxWsPortProxyFactoryBean` requires that we specify the port and service names in the `portName` and `serviceName` properties. A quick glance at the `name` attributes of the `<wsdl:port>` and `<wsdl:service>` elements in the WSDL will help you figure out what these properties should be set to.

Finally, the `namespaceUri` property specifies the namespace of the service. Among other things, the namespace will help `JaxWsPortProxyFactoryBean` locate the service definition in the WSDL. As with the port and service names, you can find the correct value for this property by looking in the WSDL. It's usually available in the `targetNamespace` attribute of the `<wsdl:definitions>` element.

10.6   Summary

Working with remote services is usually a tedious chore. But Spring provides remoting support that makes working with remote services as simple as working with any regular JavaBeans.

On the client side, Spring provides proxy factory beans that enable you to configure remote services in your Spring application. Regardless of whether you're using RMI, Hessian, Burlap, Spring's own HTTP invoker, or web services for remoting, you can wire remote services into your application as if they were POJOs. Spring even catches any `RemoteExceptions` that are thrown and rethrows runtime `RemoteAccess-Exceptions` in their place, freeing your code from having to deal with an exception that it probably can't recover from.

Even though Spring hides many of the details of remote services, making them appear as though they're local JavaBeans, you should bear in mind the consequences of remote services. Remote services, by their nature, are typically less efficient than local services. You should consider this when writing code that accesses remote services, limiting remote calls to avoid performance bottlenecks.

In this chapter, you saw how Spring can be used to expose and consume services based on some basic remoting technologies. Although these remoting options are useful in distributing applications, this was just a taste of what's involved in working within a service-oriented architecture (SOA).

We also looked at how to export beans as SOAP-based web services. Though this is an easy way to develop web services, it may not be the best choice from an architectural standpoint. In the next chapter, we'll look at a different approach to building distributed applications by exposing portions of the application as RESTful resources.